

NLF: A RESISTOR-NETWORK FRAMEWORK AND LINEAR-TIME SOLVER FOR CONVEX NETWORK-FLOW EQUILIBRIA*

OREN E. LIVNE†

Abstract. We present **NLF** (Nonlinear Laplacian Flow), a unified framework and linear-time solver for convex network-flow equilibria. A broad class of flow problems—congestion (traffic) routing, minimum-delay communication routing, and maximum flow—is the stationarity of a convex, edge-separable energy and shares one form: a *nonlinear graph Laplacian* $B\rho(B^\top\phi) = \alpha d$, in which a single monotone edge law ρ_e encodes the physics. All problems are posed in *undirected* form: flows are signed, each edge law is odd, and capacities—maximum flow included—bound the flow magnitude in either orientation; directed, arc-based variants are future work. The energy is classical—the optimality system of a convex, edge-separable network-flow program (monotropic programming); what is new is the *unification*, that one monotone edge law parametrizes all three problems, together with how it is solved: where electrical-flow and interior-point methods for flow drive an *outer* reweighting or barrier loop around *linear* Laplacian solves to reach an *approximate* optimum, NLF puts the nonlinearity in the edge law itself, so one energy encodes the problem *exactly*. NLF solves it by a damped chord-Newton iteration whose frozen global linearization—itsself a weighted graph Laplacian—is inverted by a lazily refreshed near-linear Laplacian solver applied directly to the non-linear residual; the framework is engine-agnostic (we use LAMG+, and verify approximate Cholesky is interchangeable). Each inner solve is $\mathcal{O}(m)$ and the whole nonlinear solve costs 2–4 linear Laplacian solves (up to ~ 7 on the largest social graphs). On *undirected, single-commodity* congestion over real road-network topologies (BPR cost) NLF matches a state-of-the-art interior-point method *on the identical convex program* to 10^{-10} , in a step count with no size trend, at empirically near-linear cost; on poorly-separable graphs—where the IPM’s direct KKT core is superlinear—the speed-up over that interior-point baseline widens past $45\times$ (synthetic Erdős–Rényi graphs, an extreme low-separability case). The inner solve is a replaceable module—a matrix-free first-order method on the dual ties NLF on *well-conditioned* graphs—but multigrid is the robust default: running all three paradigms (multigrid-Newton, interior-point, first-order) on the same program across 1,669 SuiteSparse graphs, *NLF converges on every one*, is fastest-or-only on 85% (median $2.6\times/4.2\times$ faster than the interior-point/first-order baselines where they converge), and is the *only* solver to finish within the time budget on the 90 hardest graphs—ill-conditioned and poorly-separable at once (first-order slowed by the former, sparse-direct by the latter). Robustness is broad: across the full SuiteSparse collection (2,003 graphs to 1.8×10^7 edges, plus giants to 5.6×10^7), under one BPR instance per graph, NLF converged on all of them at a median of 18 cycles, with a fitted time exponent $t \propto m^{0.98}$. A first multicommodity extension routes K commodities through one shared hierarchy at $\mathcal{O}(Km)$ per step—corpus-wide at $K = 4$ ($3.8\times$ single-commodity cost)—and yields a measured finding: on real capacity data, congestion couples commodities by inflating shared-corridor costs (up to $8.8\times$), not by rerouting them. The same machinery recovers the exact combinatorial max-flow as a short sequence of Laplacian solves, with the cut potential as a by-product.

Key words. network equilibrium, traffic assignment, maximum flow, algebraic multigrid, nonlinear multigrid, inexact Newton methods, continuation, graph Laplacian

AMS subject classifications. 90B20, 90C35, 65F10, 65N55, 90B10, 05C21

1. Introduction. Flows on networks at equilibrium are everywhere: drivers distributing over a road map until no route is faster (traffic assignment [9, 10, 11]), packets routed to minimize delay in a communication network [15, 16], current in a resistor network [26], and—at the combinatorial extreme—the maximum flow through a capacitated graph [17, 18, 20]. These look like different problems with different algorithms: traffic assignment is solved by Frank–Wolfe or general convex optimization, max-flow by augmenting-path or push–relabel combinatorics. We make a unifying observation and build a single fast solver on it.

*Submitted to the editors June 27, 2026.

†Pine Birch Analytics, 35 Kelinger Rd, Churchville, PA 18966-1033 (oren.livne@gmail.com, tel. 312-533-9130, pinebirchanalytics.com; ORCID: 0000-0001-6700-483X).

The observation. Each of these is the stationary point of a convex, edge-separable energy in the node potentials, and its Euler equation is one and the same *nonlinear graph Laplacian*,

$$(1.1) \quad B \rho(B^\top \phi) = \alpha d,$$

where B is the node–edge incidence matrix, ϕ the node potentials, d a source/sink demand, α a load, and ρ_e a *monotone edge law* that alone encodes the physics: a saturating ρ_e gives a hard capacity (max flow); a gently rising ρ_e gives a congestion cost (traffic). The flow is $f_e = \rho_e((B^\top \phi)_e)$, and every Newton linearization of (1.1) is an ordinary weighted graph Laplacian $J = B \text{diag}(\rho') B^\top$. A fast *near-linear* graph-Laplacian solver therefore solves all of them. What makes the *nonlinear* problem cheap is not which solver that is, but how it is used: a single global linearization—itsself a weighted graph Laplacian—is frozen across steps, lazily refreshed, and inverted directly against the nonlinear residual as a damped, inexact chord-Newton iteration, rather than solved to a tolerance of its own, with the continuation woven into the same loop; the entire nonlinear solve costs 2–4 linear Laplacian solves (§4.7). Any near-linear Laplacian solver serves as this inner engine—the framework is engine-agnostic (§3.1)—and we use **LAMG+** [8], a lean algebraic multigrid solver robust across graph classes, as a convenient parameter-free default, having verified that approximate Cholesky is fully interchangeable. One incidental convenience of LAMG+’s affinity-based aggregation is that it naturally respects the saturating min cut, though the framework does not depend on it (§3.1).

1.1. Contribution. This paper makes three contributions, in order of practical weight.

1. *A unifying resistor-network formulation* (§2) that casts maximum flow, congestion (traffic) routing, and minimum-delay routing as one nonlinear Laplacian (1.1), with the problem class fixed by a single monotone edge law and the easy/hard dichotomy (no-fold vs. fold; §2) fixed by whether that law saturates. The stationarity equation itself is classical—it is the optimality system of a convex, edge-separable network-flow program (monotropic programming); our contribution is the *unification*, that one edge law parametrizes all three problems, together with its operational consequence: a *single solver* addresses the whole class. Unlike the electrical-flow and interior-point “Laplacian paradigm” for flow [23, 24], which drive an outer reweighting/barrier loop around *linear* solves to an $(1 + \epsilon)$ -approximate optimum, here the nonlinearity sits in the edge law and the equilibrium reached is *exact*.
2. *A linear-time solver for convex congestion flow* (§4)—establishing the two claims on which any larger congestion solver rests: that the convex Beckmann equilibrium is exactly (1.1), and that a robust $\mathcal{O}(m)$ near-linear Laplacian inner solve makes it linear-time. No combinatorial algorithm addresses this problem. NLF matches a state-of-the-art IPM to 10^{-10} in a flat step count; on poorly-separable graphs NLF is the only $\mathcal{O}(m)$ solver and wins past $45\times$ (§4.4). Robustness is corpus-wide: 2,003 SuiteSparse graphs converge, $t \propto m^{0.98}$, entire nonlinear solve 2–4 linear solves (§4.5, 4.7). A first multicommodity layer (§6) carries K commodities through one hierarchy at $\mathcal{O}(Km)$; on real capacity data coupling inflates shared-corridor costs ($\leq 8.8\times$) without rerouting.
3. *An $\mathcal{O}(m)$ maximum-flow solver* (§5). The saturating instance recovers the *exact* combinatorial max-flow as a short, size-independent sequence of Laplacian solves, with the cut respecting multigrid coarsening arising for free. Not faster than a tuned combinatorial solver, but returns the cut potential and full parametric flow

98 curve.

99 Extensions—directed multicommodity assignment and a full-multigrid variant—are
100 outlined in §7.

101 **1.2. The practical problem and its applications.** Computing an equilib-
102 rium flow on a large network is a recurring industrial task, and in each domain the
103 computational bottleneck is the same Laplacian-structured solve at the core of (1.1).

104 *Road traffic (assignment).* Metropolitan planning agencies compute user-
105 equilibrium traffic assignment—how drivers distribute over a road map until no
106 route is faster—to evaluate road projects and pricing [9, 10, 11]. This is the convex
107 Beckmann program with the BPR congestion law; it is NLF’s no-fold instance and
108 the setting of our sharpest result (§4). The incumbent solvers are Frank–Wolfe and
109 bush-based methods (§4.2).

110 *Communication and data-center routing (minimum delay).* Routing packets to
111 minimize end-to-end delay under the Kleinrock $M/M/1$ model [15, 16] is the same
112 equilibrium with a capacity-saturating delay law—a fold instance handled by the
113 continuation of §3.3.

114 *Maximum flow (logistics, scheduling, vision).* The capacitated maximum-flow
115 problem [17, 18] is the saturating-law extreme of the same energy (§5); NLF recovers
116 the exact value together with the min cut and the full flow-vs-load curve. Combinato-
117 rial solvers are faster on the bare value, but do not return the smooth interior flow or
118 the parametric curve that an outer optimizer manipulates when max flow is an inner
119 step.

120 *Nonlinear resistor and analog-circuit networks.* The most literal instance: mono-
121 tone two-terminal devices (diodes, varistors, memristors) in DC steady state give
122 (1.1) with ρ_e the device I – V law—the form of DC circuit analysis and in-memory-
123 computing crossbars. As with OPF (§1.2.1) these topologies are typically structured
124 (direct’s regime); the edge would appear only on large, irregular analog networks.

125 Across all of these the operator wants not a one-off number but the equilibrium,
126 its binding-constraint structure, and its sensitivity to load, at a scale where the linear-
127 algebra cost dominates—the gap NLF targets.

128 **1.2.1. Out of scope: power-grid optimal power flow.** One natural-looking
129 application is *not* an NLF win, and we state so up front. In the widely used DC
130 approximation, optimal power flow (OPF) with transmission *line limits* is exactly
131 a capacity-constrained flow on the grid graph: the power balance $B'\theta = p$ is a
132 susceptance-weighted graph Laplacian, the line limits are the box constraints of the
133 saturating law (2.4), and the dispatch is set by the binding lines (a min cut) and their
134 shadow prices (locational marginal prices)—objects NLF returns. The mapping is ex-
135 act, and continental, contingency-constrained OPF is an active, funded need [2, 3, 4].
136 *But* transmission grids are well-separated (near-planar), the regime where sparse-
137 direct factorization of the same Laplacian is near-optimal: on the real grids we tested
138 (118 to 30,000 buses, plus the power-flow/OPF matrices embedded in SuiteSparse) di-
139 rect is *faster* than NLF by 2–5 $\times$, widening to 6 $\times$ on a 10,000-bus system (§4.6). **NLF**
140 **is therefore not a faster DC-OPF solver, and we make no such claim;** we
141 record DC-OPF as a conceptual instance, not a performance result. Full AC OPF is
142 further out of reach: non-convex, with a non-Laplacian power-flow Jacobian, outside
143 the convex monotone-edge-law class this paper addresses entirely.

144 **1.3. Prior work, and how NLF differs.** Three lines of work meet at (1.1);
145 NLF sits in the gap none of them fills.

The Laplacian paradigm for flow. Electrical-flow algorithms [23] and interior-point methods for generalized flow [24] solve a *sequence of linear* weighted-Laplacian systems inside an outer multiplicative-weights or barrier loop, reaching a $(1 + \epsilon)$ -approximate optimum; the flow is a *linear* function of the potential and all nonlinearity lives in the outer schedule. NLF instead places the nonlinearity in the constitutive law $f = \rho(B^\top \phi)$, so a single energy encodes the capacity *exactly* and the equilibrium reached is exact, not $(1 + \epsilon)$ (§2).

Fast linear Laplacian solvers. Nearly-linear solvers for $Lx = b$ —combinatorial multigrid [6], support-tree/Spielman–Teng theory [5], approximate Cholesky [25, 35], and algebraic multigrid including LAMG [1]—address the *linear, unconstrained* problem only. They are the engine, not the method: NLF calls a robust $\mathcal{O}(m)$ linear solver (LAMG+) as its inner kernel but adds the nonlinear, capacity-constrained equilibrium layer on top. AMG has also been applied to the *square, unconstrained* power-flow equations [7]; that is a different problem class (no objective, no inequality constraints) from the constrained flow program here.

Multilevel optimization. NLF is *not* multilevel optimization in the MG/OPT sense [29, 30]: it does not build coarse *optimization* models or recurse the objective. It reduces the equilibrium to a *single* nonlinear equation (1.1) and uses multigrid only as the *linear* inner solve inside a frozen-linearization chord-Newton continuation; the nonlinearity is carried by the edge law and the outer continuation, not by a coarse-grid objective.

Combinatorial and transport-specific solvers. Augmenting-path and push-relabel max-flow [18, 21] and the almost-linear-time exact algorithm [20] target directed max/min-cost flow; Frank–Wolfe [13] and bush-based assignment (Algorithm B [37], TAPAS [38]) target the smooth-cost traffic program. Each is specialized to one problem class. NLF’s contribution is orthogonal: a *single* solver, parametrized by one edge law, that spans all of them and inherits graph-class robustness from its multigrid core. We benchmark against representatives of each family where a runnable one exists (Frank–Wolfe and Ipopt [14] for congestion, Boykov–Kolmogorov for max flow) and position the rest honestly where it does not (§4.2).

2. A unified resistor-network formulation. Let B be the $n \times m$ node–edge incidence matrix of a connected graph, $\phi \in \mathbb{R}^n$ node potentials, d a demand with $\mathbf{1}^\top d = 0$, and $\alpha \geq 0$ a load. Give each edge a smooth, strictly convex *co-potential* Ψ_e and let its derivative be the edge law

$$(2.1) \quad f_e = \rho_e(g_e), \quad \rho_e = \Psi'_e, \quad g_e = (B^\top \phi)_e,$$

a monotone map from potential gradient to flow. Equilibrium is the stationary point of the convex energy $E(\phi) = \sum_e \Psi_e((B^\top \phi)_e) - \alpha d^\top \phi$, i.e. (1.1). Equivalently, in flow variables it is the convex program

$$(2.2) \quad \min_f \sum_e \Phi_e(f_e) \quad \text{s.t.} \quad Bf = \alpha d, \quad \Phi_e = \Psi_e^* \text{ (Legendre dual)},$$

whose multipliers are the potentials ϕ ; the marginal cost is $t_e = \Phi'_e = \rho_e^{-1}$. The Newton linearization of (1.1) is the weighted graph Laplacian

$$(2.3) \quad J = B \operatorname{diag}(\rho'(B^\top \phi)) B^\top, \quad w_e = \rho'_e(g_e) = \Psi''_e(g_e) = 1/t'_e(f_e) \geq 0,$$

a *solution-dependent conductance network*. Everything specific to a problem lives in the one scalar law ρ_e ; the solver below never sees anything but ρ_e, ρ'_e and J . Appendix A derives, for each application in turn, the equilibrium principle, the convex

TABLE 2.1

Three instances of the resistor-network framework (1.1). The edge law ρ_e (flow vs. potential gradient) is the inverse of the marginal cost $t_e = \Phi'_e$. “Fold” marks a feasibility limit where the conductance $\rho'_e \rightarrow 0$ and J becomes singular.

Problem	Marginal cost $t_e(f)$	Conductance ρ'_e	Fold?	Application
maximum flow [17, 19]	box $ f \leq c_e$ (soft barrier)	$\rightarrow 0$ at cut	yes	vision, matching
min-delay routing [15, 16]	$1/(c_e - f)$ (Kleinrock)	$\rightarrow 0$ at $f \rightarrow c_e$	yes	communication
congestion [9, 10, 11]	$t_e^0(1 + b(f/c_e)^p)$	$\in (0, \rho'_{\max}]$	no	road traffic

program it is equivalent to, and the reduction of its KKT system to (1.1)—in particular, where the complementarity conditions of the directed problems go.

The fold dichotomy. The single feature that governs difficulty is whether the edge law saturates. Table 2.1 lists three standard instances.

- *Hard capacity (saturating ρ , conductance $\rightarrow 0$).* The flow cannot exceed a capacity c_e , so ρ_e is bounded and $\rho'_e \rightarrow 0$ as the edge saturates. As the load rises to a *feasibility limit*, the saturating edges form a cut and J *folds*—a continuation turning (limit) point where α caps at the min cut F^* and J becomes singular. Maximum flow ($\rho_e = c_e \tanh(g/c_e)$) and minimum-delay communication routing (Kleinrock delay $t_e(f) = 1/(c_e - f)$, capacity c_e) are of this type; they need continuation through the fold (§3.3).
- *Soft congestion (rising ρ , conductance bounded).* The cost rises with flow but imposes no hard cap, so ρ_e is unbounded and ρ'_e stays in $(0, \rho'_{\max}]$, bounded away from 0: there is *no* feasibility limit and *no* fold. Traffic assignment with the BPR cost is of this type; J is uniformly well-conditioned and NLF needs no continuation (§4).

The max-flow law in detail. For the saturating instance used in §5, take

$$(2.4) \quad \rho_e(g) = \begin{cases} c_e^+ \tanh(g/c_e^+), & g \geq 0, \\ -c_e^- \tanh(g/c_e^-), & g < 0, \end{cases} \quad \Psi_e(g) = c_e^2 \log \cosh(g/c_e),$$

a soft-capacity barrier: $\Psi_e(g) = \frac{1}{2}g^2$ for $|g| \ll c_e$ (a unit resistor) and grows only linearly, $\Psi_e \rightarrow c_e|g|$, for $|g| \gg c_e$ (a saturated edge), so $\rho_e = \Psi'_e$ clamps the flow smoothly to $-c_e^- \leq f_e \leq c_e^+$. Here $w_e = 1 - (f_e/c_e)^2 \in [0, 1]$ is a pure “edge openness,” 1 for a slack edge and 0 for a saturated one, and **the forming min-cut is exactly the zero-conductance set**—which is why algebraic-multigrid aggregation, driven by a relaxation-based strong-connection measure, refuses to coarsen across it (§5). The congestion (BPR) law is given where it is used, in §4.

Relation to prior Laplacian-paradigm formulations. Solving flow through graph-Laplacian systems is the core of a productive line of work, but NLF’s mechanism differs in a way that is the source of its novelty. Electrical-flow methods [23] and their p -Laplacian descendants compute a sequence of *linear* electrical flows $f = \text{diag}(w) B^\top \phi$ —each a fixed weighted-Laplacian solve—and steer an *outer* multiplicative-weights loop to a $(1 + \epsilon)$ -approximate max-flow; interior-point methods for (generalized) flow [24] likewise wrap a Laplacian solve in a barrier path-following loop. In both, the flow is a *linear* function of the potential and all nonlinearity lives in the outer schedule. NLF instead places the nonlinearity in the *constitutive law*: $f = \rho(B^\top \phi)$ is a non-linear, saturating function of the gradient, so the single energy of (1.1) encodes the capacity *exactly* and its feasibility limit is the *exact* min cut, not a $(1 + \epsilon)$ relaxation; the outer loop is then a short continuation through one scalar fold, not a reweighting




FIG. 2.1. The three instances of the framework (one row per line of Table 2.1); column (A) the edge law ρ_e (flow), its derivative ρ'_e (the solution-dependent conductance) and the energy Ψ_e ; column (B) the behavior it induces. Row 1, max-flow (tanh): the law saturates at $\pm c_e$ and $\rho'_e \rightarrow 0$; (B1) the saturated edges are the min cut = the active set, and ϕ steps across them. Row 2, minimum-delay routing (Kleinrock $1/(c_e - f)$): the law saturates at capacity—the fold class, handled by the continuation of §3.3; (B2) the real application: the Abilene/Internet2 backbone colored by the computed NLF min-delay equilibrium for a Seattle→New York demand. Row 3, congestion (BPR): the law is unbounded and ρ'_e is bounded away from zero—no fold, pure cycling; (B3) the real application: the Sioux Falls road network colored by the computed NLF equilibrium utilization $|f_e|/c_e$ for a single source-sink demand (§4).

or barrier schedule. The Beckmann congestion instance (2.2) is itself classical [9]; what is new is the saturating-resistor reformulation of max-flow and the unified monotone-edge-law family that places max-flow, congestion, and minimum-delay routing under one nonlinear Laplacian, each solved by the same algebraic-multigrid inner solve.

3. The NLF algorithm. NLF is a *continuation loop over the load α* . Each fixed- α subproblem is solved by a handful of multigrid cycles applied *directly to the nonlinear residual*: a LAMG+ hierarchy, built once from the linearization $J = B \text{diag}(\rho') B^\top$ (2.3) at some earlier iterate and *frozen* across steps, is cycled against $-r(\phi) = \alpha d - B\rho(B^\top \phi)$, and the resulting correction is accepted through a line search. In classical terms this is a *damped, inexact chord-Newton iteration*: a global

linearization, frozen—operator and hierarchy alike—across steps and re-formed only when an observed-convergence monitor says it has drifted too far, its inverse applied by $\mathcal{O}(m)$ multigrid cycles rather than by an inner solve to a tolerance of its own. We present the algorithm from its simplest component to its most complex: the linear engine (§3.1), the fixed-load solve (§3.2), and the continuation through the fold (§3.3). Three principles organize the design.

1. *Reuse the linear engine.* Every linearization of (1.1) is an ordinary weighted graph Laplacian, so the cycle is a lean algebraic multigrid solver run on J , unchanged.
2. *Freeze the hierarchy (lazy-setup reuse).* Across a chain of nearby linearizations the AMG hierarchy is reused rather than rebuilt—an operator-reuse / lazy-setup strategy (distinct from Brandt’s frozen- τ , which freezes the FAS coarse-grid τ correction, a right-hand-side quantity [34, §15]): touch the setup rarely, and let an observed-convergence monitor decide when the frozen levels have drifted too far.
3. *Continue only through folds.* Where the law saturates, α folds at the feasibility limit; the well-conditioned continuation variable is the cut-mode amplitude. Where it does not, no continuation is needed.

3.1. The inner Laplacian solve. The outer loop requires only a near-linear solver for graph Laplacians; the framework is agnostic to which one. We use LAMG+ as a convenient, parameter-free default. LAMG+ is a lean caliber-1 aggregation algebraic multigrid solver for graph Laplacians: it aggregates nodes by a relaxation-based affinity, bounds the coarse-operator energy by a guard rather than an aggregate-size cap, eliminates low-degree nodes by exact Schur complements, and accelerates a fractional cycle by min-residual iterate recombination. The companion paper [8] establishes its two properties we rely on: it solves $L\phi = b$ in $\mathcal{O}(m \log(1/\varepsilon))$ work with *no parameter tuning* across 1,374 real-world graphs of every class, and on the large poorly-separable graphs—where sparse direct methods exhaust memory—it is among the fastest solvers available. The leading alternative near-linear graph-Laplacian solvers—approximate Cholesky [35] and Napov–Notay’s degree-aware aggregation (DRA-QC) [36]—are equally admissible. The framework is *engine-agnostic*: it requires only a near-linear Laplacian solve that does not stall, and any of these may be substituted. We have verified directly that approximate Cholesky is fully interchangeable as the inner solver, matching LAMG+’s convergence across the corpus (§4.5). The formulation and the frozen-Newton/continuation outer method are the contribution; the inner solve is a module. But the *choice* of module is not cosmetic, and we state its role honestly. Because (1.1) is a single nonlinear equation, one can even bypass the Newton inner solve entirely and minimize the convex dual energy $E(\phi) = \sum_e \Psi_e((B^\top \phi)_e) - \alpha d^\top \phi$ by a matrix-free quasi-Newton method (L-BFGS on the dual). On individual *well-conditioned* graphs this matches NLF in wall-clock—so the bare speed-up over the interior-point method on the well-conditioned poorly-separable instances of §4.4 is a *fill-in* effect any matrix-free solver inherits, not a property of multigrid. But that parity is the exception, not the rule. We ran all three paradigms—NLF (multigrid-Newton), Ipopt (interior-point / sparse-direct KKT), and L-BFGS-on-the-dual (matrix-free first-order)—on the *same* undirected BPR congestion program across the SuiteSparse corpus—the 1,669 graphs completed under the heavier three-solver protocol, of the 2,003 in the single-solver sweep (§4.5); the remainder are the largest, where the sparse-direct competitor is increasingly size-limited. Each competitor is capped at $5\times$ the measured NLF wall-clock (a 1 s floor and a 120 s ceiling; Table 3.1, Fig. 3.1). **NLF converged on every graph (100%);**

TABLE 3.1

Inner-solver robustness frontier across the SuiteSparse corpus (1,669 graphs, undirected BPR congestion—the same convex program for all). NLF (multigrid-Newton) vs. Ipopt (interior-point/sparse-direct KKT) vs. L-BFGS-on-the-dual (matrix-free first-order); each competitor capped at $5\times$ the measured NLF wall-clock (1 s floor, 120 s ceiling). NLF converges on every graph and is the only paradigm that stays within budget on both axes. Of the 92 union graphs neither competitor finishes within $5\times$ NLF on 90; 2 more also exceed the sparse-direct size limit (“within budget” is a wall-clock criterion, not a claim the others cannot converge given more time). Fig. 3.1 plots the frontier.

regime	graphs	outcome
both competitors converge (well-conditioned)	1,137	NLF $\leq$ both (median $2.6\times/4.2\times$ faster than Ipopt/L-BFGS)
ill-conditioned (first-order degrades)	389	L-BFGS exceeds $5\times$ NLF; NLF, Ipopt converge
poorly-separable (direct fills in)	51	Ipopt over budget (fill-in); NLF, L-BFGS converge
both over budget $\Rightarrow$ union	92	NLF is the only solver that finishes
total	1,669	NLF 100% converged; fastest-or-only on 85%

within that budget the first-order method did not finish on 481 of 1,669 and the interior-point method on 143—of the latter, eight reach the 120 s ceiling, two are too large to factor, and the rest ran slower than $5\times$ NLF. Where the competitors *do* converge, NLF is a median $2.6\times$ faster than Ipopt and $4.2\times$ faster than L-BFGS, and it is the fastest-or-only solver on 85% of the corpus. The two limiting modes are the two axes: a preconditioner-free first-order method slows as the conductance-weighted Laplacian ill-conditions ($\kappa(J)$ large—e.g. grids, $\kappa \sim N$, $\Theta(\sqrt{N})$ iterations), while sparse-direct fills in as separators worsen. **Multigrid is the inner default because it is the only paradigm that stays within budget on both axes:** on 90 corpus graphs neither competitor finishes within $5\times$ NLF—they sit at the hard corner of both axes (first-order to a median 200 and up to 1,037 dual iterations; sparse-direct filling in)—while NLF does. This is where NLF’s speed and robustness compound; it is not a claim that the others cannot converge given unbounded time (2 further graphs are also too large for sparse-direct; Fig. 3.1, upper right). On any single well-behaved class a simpler inner solve suffices; across the real-world corpus, multigrid is the robust choice.

What the outer loop freezes and lazily refreshes is the inner Laplacian solve itself—whether a multigrid hierarchy or a sparse approximate factorization—applied to the nonlinear residual; the choice does not affect convergence (§4.5). One convenient property of LAMG+’s relaxation-based aggregation is that it is *cut-respecting*: as the cut conductances vanish they decouple relaxation across the cut, so test vectors decorrelate and the affinity declines to aggregate across it; the measured $\mu \approx 0.01$ then holds through the continuation (§5), with no fold-aware tuning. This is a feature of the default engine, not a requirement of the framework: adequate refresh of the frozen operator near the fold suffices for *any* inner solver, and over-freezing degrades all of them equally. (α , the continuation, and the fold never enter the linear solve; the singular cut direction near F^* is deflated outside the engine, §3.3.)

3.2. The fixed-load solve: frozen-hierarchy cycles. The problem at a fixed load α is (1.1) itself: find the zero-mean potentials ϕ with $B\rho(B^\top\phi) = \alpha d$, i.e. drive the nonlinear residual $r(\phi) = \alpha d - B\rho(B^\top\phi)$ to zero. NLF solves it by Algorithm 3.1. Each pass of the loop costs one $\mathcal{O}(m)$ residual evaluation and one $\mathcal{O}(m)$ LAMG+ cycle (run inexactly, to the forcing fraction $\eta\|r\|$, $\eta = 0.05$ [32]); the expensive setup is paid only on refresh. The cycle applies an approximate *inverse* of the frozen linearization, $\delta \approx J_H^{-1}r$ —the (frozen-)Newton direction computed by multigrid. With a fresh hierarchy




FIG. 3.1. *Inner-solver robustness frontier over 1,669 SuiteSparse graphs (undirected BPR congestion): each point is a graph, placed by a conditioning proxy (x : L-BFGS-on-dual iterations) and a separability proxy (y : Ipopt time / NLF time), colored by outcome. First-order (L-BFGS) runs over budget toward the right (ill-conditioned); sparse-direct (Ipopt) toward the top (poorly-separable); in the upper right both exceed the $5 \times$ NLF budget and NLF/multigrid is the only solver that finishes. NLF converged on all 1,669.*

Algorithm 3.1 NLF fixed-load solve

Require: load α ; warm start ϕ ; hierarchy state H (shared across calls); relative residual tolerance tol (default 10^{-9}); forcing parameter η (default 0.05)

- 1: $r \leftarrow \alpha d - B\rho(B^\top \phi)$
 - 2: **while** $\|r\| > \text{tol} \cdot \alpha \|d\|$ **do**
 - 3: **refresh check:** if H is empty, or the previous pass reduced $\|r\|$ by a factor worse than 0.25, or step 5 stalled on a stale H , rebuild $H \leftarrow \text{LAMG+ setup of } J(\phi) = B \text{diag}(\rho'(B^\top \phi))B^\top$
 - 4: $\delta \leftarrow$ approximate solution of the frozen correction equation $J_H \delta = r$, where $J_H = B \text{diag}(\rho'(B^\top \phi_{\text{build}}))B^\top$ is the linearization H was built from: LAMG+ cycle(s) on H starting from $\delta = 0$, stopped when $\|r - J_H \delta\| \leq \eta \|r\|$
 - 5: **line search:** for $\tau = 1, \frac{1}{2}, \frac{1}{4}, \dots$ accept the first τ with $\|\alpha d - B\rho(B^\top(\phi + \tau\delta))\| \leq \|r\|$ and set $\phi \leftarrow \phi + \tau\delta$; on failure with a stale H , refresh and retry once
 - 6: $r \leftarrow \alpha d - B\rho(B^\top \phi)$
 - 7: **end while**
 - 8: **return** ϕ
-

and an exact solve the step would be exactly Newton's; frozen and inexact, it is a chord-Newton iteration whose linear rate the refresh monitor keeps below 0.25.

The damped update. The line search (step 5) is standard backtracking on the residual norm [31, Ch. 3]. Because δ comes from a frozen linearization, a full step can overshoot right after a load increase where the law's curvature is strongest; backtracking by halving and accepting the first non-increasing residual guarantees monotone progress at negligible cost. Near the solution the full step is always accepted ($\tau = 1$,

one trial, measured), so asymptotically the damping is inactive.

When the hierarchy is refreshed. Three triggers, all observed-convergence based, no tuning: (i) no hierarchy yet; (ii) the per-step residual factor exceeded 0.25—the frozen J has drifted too far from the current linearization; (iii) the line search stalled on a stale hierarchy (refresh, retry once; only a fresh-hierarchy stall terminates). Measured across the corpus of §4.5, a handful of refreshes per solve suffice (often one).

Deep hierarchies. On graphs of huge diameter (≥ 20 -level hierarchies: country-scale roads, large FEM meshes) the source-sink dipole right-hand side maximally excites the longest-wavelength mode, which the stock cycle schedule under-treats (factor ≈ 0.9). Such hierarchies run a *grown* cycle index ($1.15\times$ per level, per-level cap $\gamma\pi \leq 0.95$, so each cycle remains $\mathcal{O}(m)$); if the factor exceeds 0.6, the schedule is regrown and the solve restarted—restart is essential because a slow phase leaves the error in a subspace no admissible schedule contracts. This engages on 11 of 2,003 corpus graphs (0.5%); on the affected class it restores textbook behavior (**roadNet-PA**: 100-cycle stall $\rightarrow$ 5 cycles, 446 s $\rightarrow$ 20 s).

Relation to full global linearization. Rebuilding the hierarchy every step computes the same iterates at $\mathcal{O}(m)$ complexity but pays the setup at each step: measured across network and graph classes, $1.5\text{--}4.7\times$ slower wall-clock (**refresh** = 0 enables this). The one-shot FAS alternative is examined and sealed in §7.2.

3.3. Continuation through the fold (saturating laws only). For the no-fold congestion class, the whole solver is Algorithm 3.1 inside a short load ramp (Algorithm 3.2, left): the ramp is plain warm-starting, needed only to globalize the first iterations, and the hierarchy state H is shared across the entire chain.

A saturating law is different: a fixed- α solve is feasible iff $\alpha < F^*$ (the min-cut value), and as $\alpha \uparrow F^*$ the solution curve *folds*— $\alpha(V)$ flattens onto the asymptote F^* while the potentials grow without bound (Fig. 3.2); at the fold $\kappa(J) \sim (F^* - \alpha)^{-1}$ and the source-sink gap is $V = |\phi_s - \phi_t| \sim -\log(F^* - \alpha)$. Stepping α itself is therefore the worst choice: the feasible loads crowd at the fold and the step is forced to zero. The well-conditioned variable is the *cut-mode amplitude* $\psi = \chi^\top \phi \sim -\log(F^* - \alpha)$,¹ χ the (a priori unknown) min-cut indicator. We realize this by standard **pseudo-arclength continuation** [27, 28] (Algorithm 3.2, right), which never forms χ : the novelty is not the continuation but *what folds*—a graph-Laplacian whose singular direction is the min-cut indicator—and that the inner solve is $\mathcal{O}(m)$ multigrid. The solution-curve tangent satisfies $J\dot{\phi} = \dot{\alpha}d$,

$$(3.1) \quad \dot{\phi} = J^+d \quad (\dot{\alpha} = 1), \text{ then normalize } (\dot{\phi}, \dot{\alpha}),$$

one linear cycle on the (freshly refreshed) hierarchy; as $\alpha \rightarrow F^*$, $J^+d \propto \chi/\lambda_{\min}$, so the tangent rotates to align with the cut mode automatically. Each predictor along the tangent is corrected by iterating on the bordered system

$$(3.2) \quad \begin{bmatrix} J & -d \\ \dot{\phi}^\top & \dot{\alpha} \end{bmatrix} \begin{bmatrix} \delta\phi \\ \delta\alpha \end{bmatrix} = - \begin{bmatrix} r \\ N \end{bmatrix},$$

¹Justification of both asymptotics. Near the fold $\phi = V\chi + O(1)$, χ the source-side indicator, so every cut edge carries gradient $g_e = V + O(1)$. The saturating tail of (2.4) is $\rho_e(g) = c_e \tanh(g/c_e) = c_e - 2c_e e^{-2g/c_e} + O(e^{-4g/c_e})$, and all s - t flow crosses the cut, so $\alpha = \sum_{e \in \text{cut}} \rho_e(g_e) = F^* - 2 \sum_{e \in \text{cut}} c_e e^{-2V/c_e} (1 + o(1))$. Hence $F^* - \alpha \asymp e^{-2V/\bar{c}}$, $\bar{c}$ the largest cut capacity, i.e. $V = \frac{\bar{c}}{2} \log \frac{1}{F^* - \alpha} + O(1)$ and $\psi = \chi^\top \phi \propto V$. The same expansion gives the cut conductances $\rho'_e(g_e) = \text{sech}^2(g_e/c_e) = 4e^{-2V/c_e} (1 + o(1))$, whence $\lambda_{\min}(J) \asymp \sum_{\text{cut}} \rho'_e \propto F^* - \alpha$ (Rayleigh quotient at χ , which loads only the cut edges) and $\kappa(J) \sim (F^* - \alpha)^{-1}$.

Algorithm 3.2 NLF continuation. Left: no-fold (congestion). Right: fold (max-flow).

- 1: **No fold:** for $\ell = \frac{1}{4}, \frac{1}{2}, 1$: run Algorithm 3.1 at load $\ell\alpha$, warm-started, H shared.
 return ϕ
 - 2: **Fold:** climb α geometrically (each load by Algorithm 3.1) until infeasible—this brackets F^*
 - 3: refresh H ; tangent $(\dot{\phi}, \dot{\alpha})$ by (3.1)
 - 4: predictor: $(\phi, \alpha) += \Delta s(\dot{\phi}, \dot{\alpha})$; corrector: iterate (3.2) with line search, H frozen (refresh on factor > 0.25 or stall; a corrector that fails on a *stale* H is retried once at the same Δs with a fresh one—only a fresh- H failure halves Δs)
 - 5: on success grow Δs ; stop when $\dot{\alpha} \rightarrow 0$ (the fold): $\alpha = F^*$. Else go to 3
-

368 N the arclength residual. Block-eliminating $\delta\alpha$ shows what the engine actually inverts:
 369 with $u = J^+d$ and $w = J^+r$ (two applications of the inner solver),

$$370 \quad (3.3) \quad \delta\alpha = \frac{\dot{\phi}^\top w - N}{\dot{\alpha} + \dot{\phi}^\top u}, \quad \delta\phi = \delta\alpha u - w.$$

371 Both $u, w \in \text{range}(J^+) = \chi^\perp$: the near-null cut direction χ is pinned by the scalar
 372 arclength row, not by these solves, so the cycles resolve J only on the *deflated comple-*
 373 *ment* $\chi^\perp$, at the measured uniform $\mu \approx 0.01$. The border keeps the bordered system
 374 nonsingular through the fold ($\kappa = \mathcal{O}(10^3)$ while $\kappa(J) \rightarrow \infty$).

375 The refresh policy of §3.2 is woven into the corrector as a *staleness safeguard*
 376 (Algorithm 3.2, step 4). The conductances drain exponentially in the cut amplitude,
 377 so the frozen hierarchy ages faster as the fold sharpens—measured, the monitor re-
 378 freshes about every 4 steps at $0.97F^*$ and every 2 at $0.995F^*$, automatically; the
 379 tangent solve, the direction most sensitive to J 's drift, always uses a fresh hierarchy
 380 (once per accepted step). With these, the solver converges to machine precision at
 381 fixed loads up to $0.995F^*$ (relative residual $\sim 10^{-14}$) and tracks the tight-tolerance
 382 continuation without stalling.

383 4. Undirected congestion equilibrium on real road-network graphs.

384 This is NLF's sharpest result: on convex congestion flow—which has no combinatorial
 385 competitor and whose operator never folds—NLF is an empirically linear-time solver,
 386 cross-checked against a state-of-the-art interior-point method on real road-network
 387 topologies.

388 *Scope: an undirected, single-commodity study.* Throughout this section the
 389 congestion problem is the *undirected, single-commodity Beckmann relaxation*: flows
 390 are signed, the cost depends on $|f_e|$, and Wardrop's path-nonnegativity is dropped
 391 (App. A). This is *not* the transportation field's directed origin–destination user
 392 equilibrium, which loads opposing arcs independently; on benchmarks such as Sioux-
 393 Falls the directed and undirected objectives differ substantially. Our experiments
 394 establish two claims: (i) this undirected congestion equilibrium is exactly (1.1) for
 395 the inverse-BPR law, and (ii) a robust $\mathcal{O}(m)$ algebraic-multigrid inner solve makes
 396 it linear-time. The real road networks supply real topologies, capacities, and BPR
 397 data exercised with a single source–sink demand; we cross-check NLF against an
 398 independent high-accuracy solver *on the identical convex program*. The directed
 399 origin–destination assignment is future work (§7); §6 builds its first rung.




FIG. 3.2. *The α - V fold (max-flow).* As the load α is driven up, the source-sink gap $V = |\phi_s - \phi_t|$ grows; α rises, then folds onto the min-cut value F^* as the cut saturates ($\rho' \rightarrow 0$, J singular) with $V \sim -\log(F^* - \alpha)$. Stepping α crowds the fold; continuing in $V \approx \psi$ keeps steps uniform. Congestion laws have no such fold.

4.1. The BPR congestion law. On a congested road the travel time rises with the flow it carries. The field-standard model is the Bureau of Public Roads (BPR) cost [10],

$$(4.1) \quad t_e(f) = t_e^0 \left(1 + b(f/c_e)^p \right), \quad b = 0.15, \quad p = 4,$$

with t_e^0 the free-flow travel time and c_e the practical capacity. The undirected Beckmann congestion equilibrium (signed flows; Wardrop's path-nonnegativity relaxed, App. A) is the convex program (2.2) with $\Phi_e(f) = \int_0^f t_e = t_e^0 f + \frac{b t_e^0}{p+1} f^{p+1}/c_e^p$; its edge law is the inverse marginal cost $\rho_e = t_e^{-1}$, with conductance $\rho'_e = 1/t'_e(f) = (t_e^0 b p (f/c_e)^{p-1}/c_e)^{-1}$ that is *bounded away from zero*: the cost rises but never caps the flow, so there is no feasibility limit and no fold. The Jacobian $J = B \text{diag}(\rho') B^\top$ is a uniformly well-conditioned weighted Laplacian at every iterate, and NLF runs as pure fixed-load cycling (Algorithm 3.1)—each step an $\mathcal{O}(m)$ LAMG+ cycle, a handful of steps in total, $\mathcal{O}(m)$ overall.²

4.2. No combinatorial competitor; an interior-point baseline. Convex congestion flow is not a max-flow: augmenting-path and push-relabel do not apply, and no combinatorial algorithm solves (2.2). The field's classical workhorse, Frank-Wolfe on the traffic-assignment program [13], converges only sublinearly: on our Beckmann grid instances it needs an iteration count that grows with size to reach even a 10^{-4} relative duality gap (115 iterations at $n=144$) and cannot reach the high accuracy NLF attains—a textbook limitation of its $\mathcal{O}(1/k)$ rate. The transportation-specific high-accuracy state of the art is the *bush-based* family (Bar-Gera's Algorithm B [37] and TAPAS [38]); we sought a runnable open-source implementation for a same-instance head-to-head and could not obtain one—the maintained package AequilibraE

²We use a strictly convex, twice-differentiable form of (4.1) with the same free-flow and degree- p congestion structure and real (t_e^0, c_e, b, p) from the data; this regularizes the free-flow threshold so the dual problem is well posed. “Practical capacity” sentinels coded as uncapacitated links are given a large finite capacity.

TABLE 4.1

Frank–Wolfe on the convex Beckmann program (synthetic Sioux-Falls-like grids): iterations to a 10^{-4} duality gap grow with size and the $\mathcal{O}(1/k)$ rate cannot reach high accuracy. NLF reaches 10^{-10} in a flat 9 steps regardless of size (Table 4.3).

n	m	FW iters	NLF steps
16	48	4	9
36	120	4	9
64	224	7	9
100	360	9	9
144	528	115	9

failed to build in our environment, and a minimal reference Algorithm B does not reach high accuracy—so we position the bush-based methods qualitatively and benchmark quantitatively against the strong, general high-accuracy baseline, **Ipopt** [14], a mature interior-point solver run on the identical program. This is conservative for NLF: Ipopt reaches the same 10^{-10} accuracy NLF does, whereas Frank–Wolfe and bush-based methods are specialized to the *smooth-cost* program and—unlike NLF—do not address the saturating, linear-objective max-flow instance at all (§5). Ipopt’s per-iteration cost is a sparse-direct factorization of the KKT system: cheap when the graph has good vertex separators, but its fill-in grows superlinearly when it does not—exactly the axis on which NLF’s multigrid core is immune.

To make the Frank–Wolfe comparison concrete we ran it on the convex Beckmann program over synthetic Sioux-Falls-like grids (Table 4.1): the iteration count to a 10^{-4} duality gap is flat on tiny instances but grows sharply with size (115 iterations at $n=144$), and the $\mathcal{O}(1/k)$ rate cannot push the gap below $\sim 10^{-4}$ at all. NLF reaches 10^{-10} in a flat 9 continuation steps, independent of size (Table 4.3). A minimal 250-line bush-based Algorithm B reference, run on the same instances, did not reach high accuracy (gaps 10^{-2} – 10^{-1}); the production bush-based and TAPAS solvers do, but no runnable open build was available for a same-instance head-to-head, so we benchmark Ipopt quantitatively and position the bush-based family qualitatively.

4.3. Real road data: NLF and Ipopt agree on the same undirected program. We use the *Transportation Networks for Research* benchmark [12], the field-standard repository of real metropolitan road-network topologies with real capacities, free-flow times and BPR parameters. Across instances spanning three decades of size, NLF and Ipopt reach the *identical* equilibrium of the undirected single-commodity Beckmann program (Table 4.2): the Beckmann objective agrees to the solver tolerance and the equilibrium link flows to $\|f_{\text{NLF}} - f_{\text{Ipopt}}\|/\|f_{\text{Ipopt}}\| \sim 10^{-10}$, in a continuation-step count of ≈ 9 with no size trend. On this undirected congestion relaxation NLF reproduces the interior-point optimum to full accuracy—an exact same-program cross-check, not a claim on the directed origin–destination user equilibrium (§4, App. A).

4.4. Linear scaling, and the direct-vs-iterative crossover. Figure 4.1 and Table 4.3 time both solvers on the real road networks (planar, good separators) and on synthetic Erdős–Rényi graphs carrying BPR parameters (poorly separable), to 2.4×10^5 edges. NLF’s per-edge cost is flat on *both* families—empirically near-linear (the full-corpus fit of §4.5 gives $t \propto m^{0.98}$); neither the step count (≈ 9 , no size trend) nor the multigrid factor sees the graph’s separability. Ipopt does. On the real road networks its sparse-direct core is near-linear and *beats* NLF (by $\approx 3\times$; planar graphs are where direct methods excel). On poorly-separable graphs the KKT fill-in

TABLE 4.2

NLF (BPR congestion, LAMG+ inner) vs. Ipopt [12] on the undirected, single-commodity Beckmann congestion program over real road-network topologies. Both solvers reach the same equilibrium of this program; the objective matches to the tolerance, link flows to $\sim 10^{-10}$. This is the undirected net-flow relaxation, not the field’s directed origin–destination user equilibrium (§4)—it is an exact cross-check of two solvers on one identical convex program, so its objective value is not comparable to published directed-UE results for these networks. NLF’s step count shows no size trend. (Two further instances, Philadelphia and Berlin-Center, also match to 10^{-10} but need more steps under extreme capacity heterogeneity; omitted.)

network	n	m	NLF obj.	$\ \Delta f\ /\ f\ $	steps
Sioux Falls	24	38	784.96	$3 \cdot 10^{-9}$	7
Anaheim	416	634	1600.92	$6 \cdot 10^{-12}$	9
Chicago-Sketch	933	1 475	8178.62	$3 \cdot 10^{-11}$	9
Austin	7 388	10 591	16615.74	$4 \cdot 10^{-12}$	10
Chicago-regional	12 979	20 627	281.15	$2 \cdot 10^{-11}$	9
Sydney	32 956	38 787	1198.98	$2 \cdot 10^{-11}$	9

TABLE 4.3

Scaling on poorly-separable (Erdős–Rényi) graphs with BPR cost: NLF (Algorithm 3.2, $\mathcal{O}(m)$ LAMG+ inner) vs. Ipopt 3.14.19 (interior-point, MUMPS 5.9.0 sparse-direct KKT). Both reach the same optimum of the undirected congestion program. NLF’s step count shows no size trend; Ipopt is superlinear and exceeds the 120 s CPU budget on the largest instance (“—”). Speed-up = $t_{\text{Ipopt}}/t_{\text{NLF}}$.

m	NLF (s)	steps	Ipopt (s)	speed-up
12 175	0.060	9	0.152	$2.5\times$
29 939	0.173	9	0.856	$4.9\times$
59 732	0.573	9	4.97	$8.7\times$
119 460	1.30	9	58.2	$45\times$
239 842	3.01	9	— (> 120 s)	—

explodes— $t \propto m^{2.2}$ —so the gap inverts and widens with size: NLF is $5\times$ faster at $m = 3.0 \times 10^4$, $45\times$ at $m = 1.2 \times 10^5$, and at $m = 2.4 \times 10^5$ Ipopt does not finish within its 120 s CPU budget while NLF finishes in 3.0 s. This is the textbook direct-vs-iterative crossover—now for a *nonlinear* convex flow. We are careful about what it shows: these Erdős–Rényi graphs are *well-conditioned*, so the win over the IPM is its KKT *fill-in*, and *any* matrix-free solver inherits it—a first-order L-BFGS on the dual ties NLF here (Table 3.1). The advantage at scale on this class is “avoid factorization,” not “use multigrid”; multigrid’s distinctive value appears on ill-conditioned graphs (Table 3.1, §3.1).

4.5. Robustness across the graph corpus. The road networks above exercise NLF on real flow topologies; we now show it inherits LAMG+’s *graph-class* robustness in full. We impose a single-commodity BPR congestion instance (random capacities and free-flow times, a far source–sink demand) on *every* graph of the real-world SuiteSparse collection that LAMG+ was benchmarked on with at least ten nodes and a nonempty connected pattern—2,003 graphs, spanning structured grids, finite-element and circuit meshes, web, social, road, and citation networks, from 10^2 to 1.8×10^7 edges—and run NLF with the LAMG+ inner solve, all at stock settings.

Under this protocol—one BPR instance per graph, random capacities, a single far source–sink demand, stock settings—NLF converged on **all 2,003 graphs**, to residuals




FIG. 4.1. Congestion (BPR) equilibrium, NLF vs. Ipopt, on real road networks (circles) and synthetic poorly-separable graphs (squares). (A) Per-edge time: NLF is flat (near- $\mathcal{O}(m)$) on both families; Ipopt is flat on roads but rises steeply on random graphs ($\propto m^{1.2}$ per edge). (B) Total wall-clock: Ipopt wins on planar roads (cheap separators), is overtaken on random graphs by a margin that widens to 45 $\times$, and fails (time limit) at $m = 2.4 \times 10^5$ where NLF finishes in 3.0 s.

at or below the 10^{-9} tolerance, with a cycle count of *median* 18 and no size trend (each step one inexact V-cycle at forcing 0.05; range 6–33; Fig. 4.2B). We prove no convergence guarantee and do not rule out adversarial capacity distributions; the claim is empirical, on this benchmark. The total wall-clock fits $t \propto m^{0.98}$ (a fitted exponent over five decades, consistent with the inner solver’s $\mathcal{O}(m \log(1/\varepsilon))$ cost—empirically near-linear, not a proof of $\mathcal{O}(m)$) over the 1,788 graphs above 10^3 edges (Fig. 4.2A), with a per-edge cost flat across five decades (median 2.4 $\mu\text{s}/\text{edge}$, 95th percentile 9.3). The cycle count and the multigrid factor simply do not see the topology; the inner LAMG+ solve carries its parameter-free $\mathcal{O}(m)$ robustness over unchanged. The largest constants belong to eleven huge-diameter road and finite-element meshes (1.5 – 13.6×10^6 edges, e.g. `germany_osm`, `roadNet-*`), whose ≥ 20 -level hierarchies engage the grown cycle schedule of the §3 footnote: 8–25 $\mu\text{s}/\text{edge}$, a 4–10 $\times$ constant over the median at the same $\mathcal{O}(m)$ slope. Three giants confirm the scaling beyond the sweep’s cap, to `hollywood-2009` at 5.6×10^7 edges (10 min, nonlinear-to-linear ratio 4–7).

4.6. Direct vs. multigrid inner: the separability axis. The corpus comparison of §3.1 (Fig. 3.1) already pits the multigrid inner against sparse-direct factorization; we confirm the mechanism by swapping NLF’s *own* inner engine between LAMG+ and a sparse-direct factorization of the same Jacobian, the nonlinear solve held fixed. The determinant is graph *separability*, not size: on good-separator graphs—2D/3D meshes, circuits, roads even at 7×10^6 edges (`italy_osm`)—nested-dissection fill stays cheap and direct wins or ties; on poorly-separable social/web/citation graphs the fill explodes and multigrid is 4–52 $\times$ faster, and on `cage13` the factorization exhausts memory and aborts while multigrid solves it in 13 s. Across the band $200 \leq \text{nnz} \leq 3 \times 10^6$ the split is systematic—direct wins the well-separated $\sim 50\%$, multigrid the poorly-separable $\sim 30\%$ (by up to 52 $\times$), the rest within 1.2 $\times$; aggregation multigrid has its own weak class (planar `delaunay`, where direct wins), reported not hidden. The genuine application of the multigrid regime is communication-network routing on poorly-separable internet/AS/P2P topologies, where the multigrid inner wins outright (24.9 $\times$ on `p2p-Gnutella31`, 4.0 $\times$ on `as-22july06`). This is the separability axis of the frontier (Fig. 3.1); the conditioning axis is what defeats the




FIG. 4.2. *NLF over the full real-world SuiteSparse corpus (2,003 graphs, all classes, 10^2 – 1.8×10^7 edges, 100% converged). (A) total wall-clock vs. m (log-log): the fit is $t \propto m^{0.98}$ over five decades (slope-1 reference dashed), empirically near-linear, per-edge cost flat at $\approx 2.4 \mu\text{s}/\text{edge}$ median. (B) cycle count: a flat distribution (median 18, range 6–33).*

first-order baseline there.

4.7. Cost as a function of the desired accuracy. The dependence on the requested solution accuracy ε is logarithmic. Each inner solve to relative residual δ costs $\mathcal{O}(m \log(1/\delta))$: the empirically bounded multigrid factor μ gives $\log(1/\delta)/\log(1/\mu)$ cycles. The outer iteration is a frozen chord-Newton with inexact-Newton forcing η_k tracking the nonlinear residual $\|r_k\|$ (§3.1); its per-step cycle counts form a series dominated by the final solve (where $\|r_k\| \rightarrow \varepsilon$), the earlier looser steps adding only a bounded geometric tail, while the continuation and globalization run at a fixed, ε -independent path tolerance. Each contribution is thus a bounded multiple of a single linear solve, so the total cost to accuracy ε is, empirically,

$$(4.2) \quad \mathcal{O}(m \log(1/\varepsilon)).$$

The constant is the whole-nonlinear-solve-to-one-linear-solve ratio measured below (2–4), and it stays bounded because the frozen chord-Newton contraction stays bounded away from 1: on the no-fold congestion class directly (§4), and *through the fold* because the continuation deflates the diverging cut direction and keeps the bordered corrector uniformly conditioned (§3.3)—so the bound does not degrade at F^* . Figure 4.3 verifies it numerically: sweeping ε from 10^{-2} to 10^{-12} on one graph per class (web, social, road, finite-element), with inexact Newton at forcing $\eta = 0.1$ (one V-cycle per step), the total inner V-cycle count grows *linearly* in $\log(1/\varepsilon)$, at 0.24–0.60 cycles per decade of accuracy.

The nonlinear/linear cost ratio is $\mathcal{O}(1)$. The sharpest single statement of the method’s efficiency is Brandt’s yardstick: solving the *nonlinear* problem should cost only a small, size-independent multiple of solving *one linear problem* of the same kind [34, §8.3]. We measure exactly this: the full congestion solve (load continuation, all steps, all cycles, lazy refresh) against a single LAMG+ setup-and-solve of the linearization at $\alpha = 0$ (the linear resistor network $B \text{diag}(\rho'(0)) B^T \phi = d$) on the same graph. Across the classes the ratio is 1.9–4.4—grid 3.1, finite-element 2.5, road 1.9, social 3.8, web (1.9×10^6 edges) 4.4—flat in size and class: the entire nonlinear equilibrium costs 2–4 linear Laplacian solves.




FIG. 4.3. *Cost vs. desired accuracy: total inner V-cycles against $\log_{10}(1/\epsilon)$, one graph per class. The count is linear in $\log(1/\epsilon)$ (0.24–0.60 cycles per decade), confirming the total cost $\mathcal{O}(m \log(1/\epsilon))$ of (4.2).*

5. Maximum flow. The saturating instance (2.4) recovers the exact combinatorial maximum flow. Read $f = \rho(B^\top \phi)$ as an edge-flow vector: ρ 's range is the capacity box $|f_e| \leq c_e$, and $Bf = \alpha d$ is conservation, so any ϕ solving (1.1) is a capacity-feasible, conserved flow of value α . The recession slope of Ψ_e is c_e , so E has a finite minimizer iff $\alpha < F^*$; as $\alpha \uparrow F^*$ the min-cut edges saturate ($\rho'_e \rightarrow 0$ there only) and ϕ develops a cut-indicator step. By max-flow/min-cut duality the largest feasible α is F^* , the min-cut capacity, and ϕ is the LP dual. Because $f = \rho(B^\top \phi)$ is a *nonlinear* function of the gradient, it realizes the true max-flow; a *linear* resistor network computes only the sub-maximal electrical flow (0.02–0.20 F^* on heterogeneous capacities). The nonlinearity is essential.

Cut-respecting coarsening. As $\alpha \uparrow F^*$ the linearization J becomes singular along a single near-null direction—the min-cut indicator—so $\kappa(J) \rightarrow \infty$. We assume a *simple fold*: the min cut is unique, so the limiting near-null space is one-dimensional (a degenerate cut of multiplicity k , e.g. in symmetric unweighted graphs, would make it k -dimensional and require a rank- k block border in place of the scalar one below). The continuation of §3.3 *deflates* that one direction to a scalar outside the engine (the bordered/pseudo-arclength system); what the inner solver sees is the deflated complement (Eq. 3.3). *The deflation is the enabling property*—it is what keeps the inner operator well-conditioned through the fold, and it is a property of the framework, not of any one inner solver (swapping multigrid for a sparse-direct or another near-linear Laplacian solve reaches the identical F^* ; §5.1). Multigrid's relaxation-based affinity is a natural fit on the complement—since the forming cut is the zero-conductance set of J (§2), the affinity declines to aggregate across it, so the coarse hierarchy mirrors the cut without being told where it is—but it is helpful, not required. The measured multigrid factor *on the deflated operator* is $\mu \approx 0.01$, uniformly in α/F^* down to $0.999 F^*$ (Fig. 5.1); the uniformity is a property of the deflation, not of aggregation alone—on the undeflated J , μ degrades with $\kappa(J)$ as expected.

Exactness and scaling. NLF returns the exact max-flow (Table 5.1), checked against F^* from a linear program over free flows and from push-relabel (which agree to 10^{-4}), on every graph, to solver tolerance. The pseudo-arclength step count is flat in size (Fig. 5.1B, mean 6.7) and the per-edge cost is nearly flat, $t/m \propto m^{0.11}$ (near-

TABLE 5.1

NLF recovers the true max-flow exactly; the gradient/electrical relaxation falls far short. F^* : LP over free flows and push-relabel agree to 10^{-4} ; *washington* and *genrmf* are DIMACS-challenge generators [22].

graph	n	F^*	grad/ F^*	NLF/ F^*
grid2d/16	256	4.973	0.20	1.0000
grid3d/6	216	17.478	0.11	1.0000
washington	66	500.745	0.02	1.0000
genrmf	64	16.000	0.00	1.0000
bottleneck	40	1.000	1.00	1.0000




FIG. 5.1. Max-flow scaling (whole algorithm, total wall-clock), on 2D/3D grids and random graphs to $\sim 10^4$ edges, all exact at F^* . (A) Per-edge cost: nearly flat ($t/m \propto m^{0.11}$, near- $\mathcal{O}(m)$) at ≈ 1.4 ms/edge, the constant set by the per-step hierarchy rebuild; Boykov-Kolmogorov sits far below at ≈ 0.17 μ s/edge (calibration only). (B) The continuation is $\mathcal{O}(1)$: pseudo-arclength steps are flat in size (mean 6.7).

$\mathcal{O}(m)$), with a large constant set by per-step hierarchy rebuilds (the lazy refresh of §3 reduces this by 1.5–4.7 $\times$). *Calibration, not a contest*: Boykov-Kolmogorov, on its target vision grids, runs at ≈ 0.17 μ s/edge (Fig. 5.1A), far below NLF’s constant—**NLF is not meant to be a faster combinatorial max-flow solver**. Its value is being $\mathcal{O}(m)$, returning the cut potential (LP dual), the full flow-vs-load curve, and a smooth interior flow—the objects an outer method manipulates when max-flow is an inner step—one instance of a framework that, on the congestion class, has *no* combinatorial competitor at all.

5.1. The inner solver at the fold is interchangeable. The deflation, not the choice of inner solver, is what carries the continuation through the fold (§5); the inner Laplacian solve is therefore a replaceable module. We check this directly: holding NLF’s outer method fixed (chord-Newton with the deflated pseudo-arclength continuation to F^*), we swap only the inner solver—multigrid (LAMG+), approximate Cholesky [25], an exact sparse Cholesky factorization, and a diagonally preconditioned CG—on max-flow instances spanning FEM, structural, and circuit graphs (Table 5.2).

The result is modularity, not a winner. *Every inner solver that converges returns the identical F^** to solver tolerance, so the continuation is genuinely solver-agnostic. The distinction that matters is near-linear versus not: the exact factorization is fastest

TABLE 5.2

Inner-solver swap inside NLF’s max-flow continuation to F^* (wall-clock seconds; “stall” = did not converge within 120 steps; “OOM” = exact factorization size-guarded at 150k nodes). Every converging solver returns the identical F^* , so the continuation is solver-agnostic; the operative choice is near-linear versus not. Representative subset; the bake-off script and its outputs are in the released artifact. On `bcsstk38` the stalled iterative solvers were arrested at 37.0, far from the true $F^* = 45.6$ the direct solve reaches.

graph	class	n	F^*	LAMG+	approxChol	direct	diagPCG
<code>grid2</code>	2D FEM	3,296	3.003	11.4	0.75	0.24	12.3
<code>airfoil1</code>	2D FEM	4,253	6.695	29.3	1.30	0.40	20.6
<code>bodyy5</code>	aniso. FEM	18,589	3.325	307	9.5	2.0	214
<code>bcsstk38</code>	structural	8,032	45.61	stall	stall	4.0	stall
<code>circuit_4</code>	circuit	67,029	1.381	7.3	5.5	2.9	124
<code>G2_circuit</code>	circuit	150,102	4.030	1579	50.1	OOM	3011

where it fits but exhausts memory past $\sim 150k$ poorly-separable nodes (`G2_circuit`, `scircuit`), exactly the regime where a near-linear inner solve is mandatory, while the diagonal-PCG baseline—no coarse correction—is typically up to two orders of magnitude slower and the least scalable. Among the near-linear solvers, multigrid and approximate Cholesky are interchangeable in robustness; approximate Cholesky often carries the smaller constant, and either is a sound default—the framework does not depend on the choice. The single stiff case (`bcsstk38`) is instructive: there *all* iterative inner solves stall and only the exact factorization reaches F^* —the deflation removes the *global* cut singularity, not local stiffness, for which an accurate inner solve is still required.

6. Multicommodity congestion. §4 established the formulation and the $\mathcal{O}(m)$ solver for a single commodity. This section extends both, end to end, to K simultaneously routed commodities: a vector edge law that couples the commodities through shared congestion (§6.1), a solver that reuses Algorithm 3.1 and *one* scalar LAMG+ hierarchy for all K commodities (§6.2), validation against an exact-Newton reference (§6.3), and the same full-corpus robustness sweep as §4.5 (§6.5). The implementation is a modular layer over the single-commodity solver; nothing in the latter changes.

6.1. Formulation: a vector edge law. K commodities ship demands $\alpha d^1, \dots, \alpha d^K$ ($\mathbf{1}^\top d^k = 0$) over one network; commodity k has flow $f^k \in \mathbb{R}^m$ and potential $\phi^k \in \mathbb{R}^n$. Collect the per-edge flows into the vector $\mathbf{f}_e = (f_e^1, \dots, f_e^K) \in \mathbb{R}^K$ and the per-edge potential gradients into $\mathbf{g}_e = ((B^\top \phi^1)_e, \dots, (B^\top \phi^K)_e)$. The modeling question is how the commodities couple through shared congestion. Three natural choices: the signed total $\sum_k f_e^k$ (degenerate—collapses to single-commodity); the volume $\sum_k |f_e^k|$ (classical but nonsmooth, forces a per-edge complementarity akin to the max-flow cut); and the Euclidean magnitude $\|\mathbf{f}_e\|$ (smooth, strictly convex, reduces exactly to $K = 1$). We take the third for this feasibility study; the volume-coupled and directed models remain future work (§7.1). The program is (2.2) with the congestion cost charged to the magnitude of the commodity-flow vector,

$$(6.1) \quad \min_{f^1, \dots, f^K} \sum_e \Phi_e(\|\mathbf{f}_e\|) \quad \text{s.t.} \quad Bf^k = \alpha d^k, \quad k = 1, \dots, K,$$

and stationarity gives, per edge, the *vector law*

$$(6.2) \quad \mathbf{f}_e = \rho_e(\|\mathbf{g}_e\|) \frac{\mathbf{g}_e}{\|\mathbf{g}_e\|} :$$

the scalar law of (2.1) applied to the magnitude of the gradient vector and acting along it, so the K gradient components share one congested conductance. At $K = 1$, (6.2) is (2.1). The system $Bf^k = \alpha d^k$ with (6.2) is the stationarity of the convex energy $E(\phi^1, \dots, \phi^K) = \sum_e \Psi_e(\|\mathbf{g}_e\|) - \alpha \sum_k (d^k)^\top \phi^k$, strictly convex modulo the K per-commodity constants, so the equilibrium is unique. Physically the coupling turns on exactly where congestion does. On real capacity data at rush-hour load, commodity flows superpose to 3×10^{-4} —real networks place each edge in one of BPR’s two power-like regimes where path splits are load-independent—while the shared load inflates marginal costs by up to $8.8\times$: the congestion externality is slowdown, not rerouting. Synthetic instances with random capacities do reroute (12–19% flow deviation), so the corpus sweep (§6.5) exercises the solver where commodities genuinely interact.

The Newton linearization of (6.2) is the K -commodity block Laplacian whose (k, l) block is $B \text{diag}(c_e \delta_{kl} + (\rho'_e - c_e) u_e^k u_e^l) B^\top$: per edge, the $K \times K$ block

$$(6.3) \quad \frac{\partial \mathbf{f}_e}{\partial \mathbf{g}_e} = c_e I_K + (\rho'_e - c_e) \mathbf{u}_e \mathbf{u}_e^\top, \quad c_e = \frac{\rho_e(s_e)}{s_e}, \quad \mathbf{u}_e = \frac{\mathbf{g}_e}{s_e}, \quad s_e = \|\mathbf{g}_e\|,$$

with eigenvalues c_e (multiplicity $K - 1$, across the flow direction) and ρ'_e (along it): symmetric positive definite for any monotone law, with anisotropy ratio $c_e/\rho'_e = t'_e(F)F/t_e(F) \in [1, p]$ bounded by the BPR exponent.

6.2. Algorithm: one scalar hierarchy carries all K commodities. Algorithm 3.1 carries over with three substitutions and no new machinery. (i) *Stacked state.* The iterate is $\Phi = (\phi^1, \dots, \phi^K)$, the residual $R = (r^1, \dots, r^K)$, $r^k = \alpha d^k - Bf^k$, and the stopping test is on the stacked norm. (ii) *A scalar frozen operator.* Rather than freeze the $K \times K$ block linearization, NLF freezes a *scalar* hierarchy H on the single conductance $w_e = \sqrt{c_e \rho'_e}$ —the geometric mean of the block’s two eigenvalues—and computes the correction as K independent cycles, δ^k from H against r^k . One $\mathcal{O}(m)$ LAMG+ setup serves all K commodities and all chord steps. Per edge the preconditioned block spectrum is $\{\sqrt{c_e/\rho'_e}, \sqrt{\rho'_e/c_e}\}$, symmetric about 1 with spread $c_e/\rho'_e \leq p$ independent of K, m , and the iterate: the damped chord iteration contracts at a rate bounded by the law’s exponent alone, and the line search of Algorithm 3.1 supplies the damping that centers it. (iii) *A self-calibrating refresh monitor.* The block-diagonal frozen operator has an intrinsic rate floor set by the anisotropy—which no rebuild can beat—so the refresh monitor measures the contraction on the step following each rebuild and treats only degradation *beyond that baseline* as staleness; otherwise a heavily congested instance would rebuild every step to no effect.

The cost accounting is the point of the construction. A chord step costs *one* vector-law evaluation— m scalar inversions, the same count as a single-commodity step, since only $\|\mathbf{g}_e\|$ is inverted—plus K cycles and K residual products: $\mathcal{O}(Km)$, on top of one shared $\mathcal{O}(m)$ setup. The entire K -dependence is K right-hand sides through one hierarchy.

Two further variants were implemented: a flexible-CG inner solve on the full frozen block linearization (for when the anisotropy bites), and an exact assembled-block Newton as the validation reference. The FCG variant matches the block-diagonal chord in step count and wall-clock—the outer iteration is bound by the law’s nonlinearity, not inner accuracy—so the block-diagonal chord is the method of record.

6.3. Validation. A fifteen-check suite validates the layer: the frozen block Jacobian (6.3) matches the finite-difference derivative of the residual; at $K = 1$ both the

TABLE 6.1

Multicommodity validation on real road networks (TNTP), $K = 4$. “Exact Newton” assembles the full $Kn \times Kn$ block linearization and re-solves it at every damped step (the reference); NLF is the block-diagonal chord with one frozen scalar hierarchy. Deviation is the relative ℓ_2 distance between the two equilibria.

Network	n	m	Newton steps	NLF steps	Setups	Deviation
SiouxFalls	24	38	8	8	1	5×10^{-15}
Anaheim	416	634	9	16	1	5×10^{-10}
Chicago-Sketch	933	1475	9	19	1	3×10^{-10}
Winnipeg	1040	1595	6	19	1	7×10^{-11}

multigrid and the direct paths reproduce the single-commodity solver’s equilibrium; all three inner solvers agree with the exact assembled-block Newton reference; identical commodities receive identical potentials and a negated demand a negated potential (neither symmetry is imposed); and the joint-vs-alone coupling check above. Table 6.1 shows the reference comparison on real TNTP road networks at $K = 4$ (four well-separated demand dipoles, benchmark load $\alpha = 0.3 \text{ med}(c)$ per commodity): NLF reaches the exact-Newton equilibrium to 10^{-10} with a handful of chord steps and a *single* hierarchy setup—whereas the exact-Newton competitor must assemble and factor the full $Kn \times Kn$ block linearization at *every* step. This is the relevant high-accuracy baseline for the coupled multicommodity program (general convex solvers such as Ipopt reduce to the same assembled-block linear algebra); NLF matches its equilibrium while replacing the K -fold-larger block solve with one shared scalar $\mathcal{O}(m)$ hierarchy.

6.4. Applications across domains. We exercise the $K = 4$ solver on the multicommodity applications the formulation targets: real TNTP metropolitan road networks (topology, capacities, free-flow times, BPR parameters, and the four heaviest origin–destination pairs all from the benchmark data, jointly scaled to a loaded rush hour, peak $\|\mathbf{f}_e\|/c_e \approx 1.07$), and corpus-protocol instances from other coupled-flow domains where flows share one network—internet AS topology, a power transmission grid, national road networks. Every instance converges to 10^{-9} with a handful of chord steps and 1–3 hierarchy setups, from 38 edges to 1.5×10^6 (e.g. roadNet-PA, 1.5×10^6 edges, 26 steps, 2 setups).

6.5. Corpus robustness and the cost of K . The decisive test is the same full-corpus sweep as §4.5, repeated at $K = 4$: the identical 2,003 real-world SuiteSparse graphs and instance protocol, with four well-separated demand dipoles per graph, each at the benchmark load. **Every one of the 2,003 graphs converges** to 10^{-9} relative residual (Fig. 6.1), to $m = 1.8 \times 10^7$ edges, at a flat median of 18 chord steps (range 9–51) and a median of *one* hierarchy setup (maximum 5; more than two on only 120 of 2,003): the shared scalar hierarchy carries the coupled four-commodity system essentially as cheaply as it carries one. The fitted total cost is $t \propto m^{1.04}$ over four decades, median 9.1 μs per edge. Joined per graph against the single-commodity sweep, the $K = 4$ solve costs a median 3.8 $\times$ the single-commodity solve (interquartile range [3.0, 5.6])—the $\approx K$ per-step work with the setup amortized, as the cost accounting of §6.2 predicts.

The cost of K is measured directly by sweeping $K \in \{1, 2, 4, 8\}$ with nested demand sets on FEM and road graphs to 1.5×10^6 edges (olesnik0, srb1, roadNet-PA): the chord-step count and the hierarchy-setup count do not move, and wall-clock grows




FIG. 6.1. Multicommodity NLF ($K=4$) over the full real-world SuiteSparse corpus, mirroring the single-commodity sweep (grey, underneath). (A) Total wall-clock vs. edge count: all 2,003 graphs converge, $t \propto m^{1.04}$ over four decades. (B) Chord-step count: flat in size and class, median 18—unchanged from the single-commodity solver.

close to linearly in K ($K=8$ at $4.2\text{--}8.5\times$ the single-commodity cost)—one shared $\mathcal{O}(m)$ setup plus K cycle solves per step, $\mathcal{O}(Km)$ in total, as the construction promises.

7. Future work. Several directions follow from the framework; we outline them, are candid that none is yet a finished method, and seal one question—the one-shot FAS cycle—with a specific negative result (§7.2).

7.1. Directed multicommodity assignment. §6 routes K coupled commodities under a smooth magnitude coupling with signed flows; real traffic assignment is *directed* (nonnegative arc flows, costs coupling through the directed link flow), validated against published origin–destination equilibria (Sioux Falls, Anaheim, ...). Two steps separate §6 from that target: the volume coupling $\sum_k |f_e^k|$ forces a per-edge complementarity (only commodities whose potential drop matches the common congested cost may flow)—a free boundary of the same character as the max-flow cut, attacked by the continuation-with-safeguard machinery of §3.3; and scale in K (thousands of demand columns), made conceivable by the $\mathcal{O}(Km)$ per-step cost with one shared hierarchy (§6.2) plus standard per-origin aggregation. Benchmarking against Frank–Wolfe and Algorithm-B on the published equilibria is the deployable-solver milestone.

7.2. One-shot FAS: what we tried, and why we sealed it. The methodologically pure alternative to Algorithm 3.1 is a genuinely nonlinear cycle: nonlinear relaxation on the nodal balance, the same elimination-and-aggregation hierarchy, full nonlinear coarse equations with the fine-to-coarse τ correction [33, 34], the load advanced by continuation—no global linearization anywhere. We built this prototype for the saturating law and record its failure so it need not be rediscovered. In a controlled experiment at *fixed* load $\alpha = 0.6F^*$ (far below the fold), from the same warm start and with equal fine relaxation, full τ -corrected V-cycles on a frozen aggregation hierarchy reached the *same* residual after 20 cycles as Gauss–Seidel–Newton relaxation alone (slightly worse on a 16×16 grid): the coarse correction removed essentially none of the smooth error—it was *inert*. Two suspects were removed by experiment: staleness (a cut-aware hierarchy frozen near F^* converged no better than a cut-blind one frozen at $\phi = 0$, and rebuilding every step did not help) and where

α is updated (orthogonal; §7.3). The surviving suspect is the inter-grid transfer of the nonlinear state—near saturation ρ' varies by orders of magnitude within an aggregate, so a caliber-1 transfer of ϕ misrepresents the flux from which the coarse τ is computed—but the flux-based transfer [34, §4.6] that would test it was never completed. Two limits keep this honest: the prototype was bespoke, not the full LAMG+ machinery rebuilt nonlinearly, and the no-fold congestion law was never FAS-tested. We do not claim one-shot FAS is impossible—only that its coarse levels did nothing where we tested it, and the economics remove the incentive (the frozen-linearization solver already costs $1.9\text{--}4.4\times$ one linear solve).

7.3. Full multigrid: α at the coarsest level. Brandt’s nested-iteration ideal [34, §5.6] would update the global scalar α at the coarsest level rather than the finest, prolongating (ϕ, α) upward with per-level arclength finishing. A first version was $1.6\text{--}2.1\times$ faster on small instances but $2.2\times$ slower on a larger grid (per-level finishing dominates)—a non-uniform gain—so we keep the simple $\mathcal{O}(m)$ finest-level architecture and leave a tuned FMG variant to future work.

7.4. Optimal transport as a further instance. The same edge-separable energy expresses optimal transport in Beckmann (flow) form: minimizing a monotone edge-separable cost subject to $Bf = \mu - \nu$ is exactly (1.1)—the quadratic (H^{-1}) cost a single $\mathcal{O}(m)$ Laplacian solve, the W_1 limit ($p \rightarrow 1$) the non-smooth saturating member, a natural fold instance. The quadratically-regularized form, whose dual Hessian is an active-subgraph Laplacian solved in the literature by direct factorization [39], is a natural opening for the near-linear inner solve; a dedicated treatment is left to future work.

8. Scope and anticipated questions.

Scope, honestly stated. NLF has no combinatorial competitor on convex congestion (matches the interior-point optimum to 10^{-10} , empirically near-linear, wins to $45\times$ on poor separators, §4); wins $2\text{--}52\times$ and is often the only solver to complete on poorly-separable graph-Laplacian flow such as communication routing (§4.6, 6); is competitive but only asymptotically faster on well-separated meshes and roads, where direct factorization wins; returns the cut, dual, and parametric curve on max-flow without competing on the bare value (§5); and makes no speed claim on DC-OPF, a motivating application only (§1.2.1). The formulation (1.1) itself is classical; the contribution is the unification, the frozen-linearization $\mathcal{O}(m)$ solver, and the empirical demonstration (§1.3).

Why only a single source–sink demand? This feasibility study establishes the two load-bearing claims—that the (undirected) equilibrium is exactly (1.1), and that an $\mathcal{O}(m)$ inner solve makes it linear-time. The directed full origin–destination assignment is future work (§7; §6 is the first rung); against the directed bush-based/TAPAS solvers [37, 38], which target that directed program and not the undirected relaxation studied here, a same-instance comparison is the clearest remaining experiment.

Reproducibility. Every table is produced by scripts run on public benchmarks—the SuiteSparse collection, the Transportation Networks road data [12], and the MATPOWER/pglib-opf grids [2]. The interior-point comparator is Ipopt 3.14.19 with the MUMPS 5.9.0 sparse linear solver and a 120s CPU budget; the BPR law is regularized to the strictly convex, twice-differentiable form noted in §4.1, and the random capacity/free-flow distributions and demand-selection seeds are fixed and recorded in the released scripts. The LAMG+ linear core is publicly available [8]; the NLF nonlinear solver, the exact instance-generation scripts, and a reproducibility

notebook—which re-runs the fold, no-fold, and competitor timings live and checks them against these tables—are released at <https://github.com/orenlivne/nlf>.

9. Conclusion. NLF is a single resistor-network framework for convex, edge-separable network-flow equilibria $B\rho(B^\top\phi) = \alpha d$: placing the nonlinearity in the edge law lets one energy encode problems from electrical flow to max-flow exactly, solved by a damped chord-Newton on a frozen linearization cycled against the non-linear residual by a lazily refreshed near-linear Laplacian solver. Its sharpest result is on undirected convex congestion—an interior-point optimum reproduced to 10^{-10} at a size-independent step count, near-linear on poorly-separable graphs where direct factorization is superlinear. Direct methods win on well-separated graphs and combinatorial solvers on the bare max-flow value; the directed origin–destination assignment remains future work (the multicommodity layer is its first rung). The contribution is the unification and the demonstration that one frozen-linearization, multigrid-inner solver handles all of it at near-linear cost.

Appendix A. Why each application is an instance of the framework.

Per application: equilibrium principle $\Leftrightarrow$ convex program (2.2) $\Leftrightarrow$ KKT $\Leftrightarrow$ $B\rho(B^\top\phi) = \alpha d$. Each instance differs only in the edge law; directed complementarity is absorbed into that law or relaxed by the signed formulation.

The chain in general. Stationarity of (2.2) in f_e gives the nonlinear Ohm law $t_e(f_e) = (B^\top\phi)_e$; inversion gives $f_e = \rho_e((B^\top\phi)_e)$; and Kirchhoff ($Bf = \alpha d$) yields (1.1). Summing along any source–sink path telescopes: every path has marginal cost $\phi_s - \phi_t$, the scalar equilibrium cost. The signed formulation has no inequality constraints—a smooth equation—which is why Newton with a near-linear Laplacian solve applies with no active sets.

Electrical flow (linear law). Ohm’s law is linear, $f_e = g_e/r_e$; the energy is the dissipated power and (2.2) is Thomson’s principle (current minimizes dissipation) [26]. Equation (1.1) is the ordinary weighted Laplacian system—the framework’s fixed point under linearization, and the object the inner solver [8] is built for.

Congestion (traffic) equilibrium. Wardrop’s principle—per origin–destination pair, path flows $h_p \geq 0$ with $\sum_p h_p = d$ obeying $T_p - \pi \geq 0$ and $h_p(T_p - \pi) = 0$ —is, by Beckmann’s theorem [9], the KKT system of $\min_h \sum_e \int_0^{f_e} t_e$: differentiating gives $\partial F / \partial h_p = \sum_{e \in p} t_e = T_p$, and π (the multiplier of $\sum_p h_p = d$) is both the marginal cost of demand and the equilibrium travel time. In link variables the demand multiplier unfolds into the node field ϕ ($\pi = \phi_s - \phi_t$) and the program becomes (2.2) with $\Phi'_e = t_e$, the BPR latency of §4. The signed relaxation keeps the first Wardrop clause exactly (all routes at potential-drop cost π) and discards $h_p \geq 0$: an unused route may carry counterflow—the undirected scope, with per-commodity nonnegativity the free-boundary extension of §7.

Minimum-delay routing. M/M/1 queueing with Little’s law gives the per-packet link delay $t_e(f) = 1/(c_e - f)$, and Kleinrock’s independence approximation makes total delay edge-separable [15]. Taking the delay as marginal cost, $\Phi_e = \log(c_e/(c_e - f))$, gives (2.2) directly; the inverse law $\rho_e(g) = c_e - 1/g$ (odd-extended) saturates at capacity with conductance $\rho'_e = (c_e - f)^2 \rightarrow 0$, placing the problem in the fold class of Table 2.1: the capacity constraint $f < c_e$ is never imposed—the law cannot produce it—and the feasibility limit, where the filling edges form a cut in the spare capacities, is the fold handled by §3.3.

Maximum flow. The LP $\max \alpha$ s.t. $Bf = \alpha d$, $|f_e| \leq c_e$ has the box constraint as its only nonsmoothness; its complementarity (an edge is either slack with zero reduced

cost or saturated on the cut) is absorbed into the smooth saturating law (2.4), whose energy $\Psi_e = c_e^2 \log \cosh(g/c_e)$ is quadratic for slack edges and exactly linear, $c_e|g|$, for saturated ones. The active set is not tracked: it *emerges* as the zero-conductance set $w_e \rightarrow 0$, the forming min cut, and the LP's degeneracy reappears as the one scalar fold at $\alpha = F^*$ traversed by the continuation of §3.3. At the fold the saturated set is the exact min cut and ϕ steps across it (§5); no $(1 + \epsilon)$ relaxation is involved.

REFERENCES

- [1] O. E. Livne and A. Brandt, *Lean Algebraic Multigrid (LAMG): Fast Graph Laplacian Linear Solver*, SIAM J. Sci. Comput. **34**(4), B499–B522, 2012.
- [2] R. D. Zimmerman, C. E. Murillo-Sánchez, and R. J. Thomas, *MATPOWER: Steady-state operations, planning, and analysis tools for power systems research and education*, IEEE Trans. Power Syst. **26**(1), 12–19, 2011.
- [3] C. Coffrin, R. Bent, K. Sundar, Y. Ng, and M. Lubin, *PowerModels.jl: an open-source framework for exploring power flow formulations*, Power Systems Computation Conf. (PSCC), 2018.
- [4] U.S. DOE ARPA-E, *Grid Optimization (GO) Competition*, 2019–2023; <https://gocompetition.energy.gov>.
- [5] D. A. Spielman and S.-H. Teng, *Nearly-linear time algorithms for graph partitioning, graph sparsification, and solving linear systems*, SIAM J. Comput. **40**(4), 981–1025, 2011.
- [6] I. Koutis, G. L. Miller, and D. Tolliver, *Combinatorial preconditioners and multilevel solvers for problems in computer vision and image processing*, Comput. Vis. Image Underst. **115**(12), 1638–1646, 2011.
- [7] K. Lee, H. Wang, and D. Bindel, *A nonlinear algebraic multigrid framework for the power flow equations*, SIAM J. Sci. Comput. **41**(3), B812–B833, 2019.
- [8] O. E. Livne, *LAMG+: A Robust Lean Algebraic Multigrid Solver for Graph Laplacians*, arXiv:2606.24791, 2026; code at <https://github.com/orenlivne/lamgplus>.
- [9] M. Beckmann, C. B. McGuire, and C. B. Winsten, *Studies in the Economics of Transportation*, Yale Univ. Press, 1956 (traffic equilibrium).
- [10] Bureau of Public Roads, *Traffic Assignment Manual*, U.S. Dept. of Commerce, Urban Planning Division, 1964.
- [11] S. D. Boyles, N. E. Lownes, and A. Unnikrishnan, *Transportation Network Analysis, Volume I: Static and Dynamic Traffic Assignment*, 2nd ed., open textbook, sboyles.github.io/blubook.html, 2023 (modern treatment of Beckmann–Wardrop equilibrium and its algorithms).
- [12] Transportation Networks for Research Core Team, *Transportation Networks for Research*, github.com/bstabler/TransportationNetworks (accessed 2026).
- [13] M. Frank and P. Wolfe, *An algorithm for quadratic programming*, Naval Res. Logist. Quart. **3**(1–2), 1956.
- [14] A. Wächter and L. T. Biegler, *On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming*, Math. Program. **106**(1), 2006 (Ipopt).
- [15] D. Bertsekas and R. Gallager, *Data Networks*, 2nd ed., Prentice Hall, 1992 (minimum-delay routing, Kleinrock cost).
- [16] A. Mendiola, J. Astorga, E. Jacob, and M. Higuero, *A survey on the contributions of software-defined networking to traffic engineering*, IEEE Commun. Surveys Tuts. **19**(2), 918–953, 2017 (modern survey of delay/utilization-driven routing optimization).
- [17] R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, *Network Flows: Theory, Algorithms, and Applications*, Prentice Hall, 1993.
- [18] A. V. Goldberg and R. E. Tarjan, *A new approach to the maximum-flow problem*, J.

- ACM **35**(4), 1988 (push-relabel).
- [19] A. V. Goldberg and R. E. Tarjan, *Efficient maximum flow algorithms*, Commun. ACM **57**(8), 82–89, 2014 (survey).
- [20] L. Chen, R. Kyng, Y. P. Liu, R. Peng, M. Probst Gutenberg, and S. Sachdeva, *Maximum flow and minimum-cost flow in almost-linear time*, Proc. 63rd IEEE Symp. Foundations of Computer Science (FOCS), 612–623, 2022.
- [21] Y. Boykov and V. Kolmogorov, *An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision*, IEEE TPAMI **26**(9), 2004.
- [22] D. S. Johnson and C. C. McGeoch (eds.), *Network Flows and Matching: First DIMACS Implementation Challenge*, AMS, 1993 (`genrmf`, `washington`).
- [23] P. Christiano, J. A. Kelner, A. Madry, D. A. Spielman, and S.-H. Teng, *Electrical flows, Laplacian systems, and faster approximate maximum flow*, STOC, 2011.
- [24] S. I. Daich and D. A. Spielman, *Faster approximate lossy generalized flow via interior point algorithms*, STOC, 2008.
- [25] R. Kyng and S. Sachdeva, *Approximate Gaussian elimination for Laplacians—fast, sparse, and simple*, FOCS, 2016 (`approxChol`, implemented in `Laplacians.jl`).
- [26] P. G. Doyle and J. L. Snell, *Random Walks and Electric Networks*, Mathematical Association of America, 1984.
- [27] H. B. Keller, *Numerical solution of bifurcation and nonlinear eigenvalue problems*, in *Applications of Bifurcation Theory*, Academic Press, 1977 (pseudo-arclength).
- [28] E. L. Allgower and K. Georg, *Introduction to Numerical Continuation Methods*, SIAM Classics in Applied Mathematics **45**, 2003.
- [29] S. G. Nash, *A multigrid approach to discretized optimization problems*, Optim. Methods Softw. **14**(1–2), 99–116, 2000.
- [30] S. Gratton, A. Sartenaer, and Ph. L. Toint, *Recursive trust-region methods for multiscale nonlinear optimization*, SIAM J. Optim. **19**(1), 414–444, 2008.
- [31] J. Nocedal and S. J. Wright, *Numerical Optimization*, 2nd ed., Springer, 2006.
- [32] R. S. Dembo, S. C. Eisenstat, and T. Steihaug, *Inexact Newton methods*, SIAM J. Numer. Anal. **19**(2), 400–408, 1982.
- [33] A. Brandt, *Multi-level adaptive solutions to boundary-value problems*, Math. Comp. **31**(138), 1977 (Full Approximation Scheme).
- [34] A. Brandt and O. E. Livne, *Multigrid Techniques: 1984 Guide with Applications to Fluid Dynamics, Revised Edition*, SIAM, 2011 (frozen- τ §15; discontinuities §4.6; global constraints/continuation §5.6, §8.3).
- [35] Y. Gao, R. Kyng, and D. A. Spielman, *Robust and practical solution of Laplacian equations by approximate elimination*, arXiv:2303.00709, 2023.
- [36] A. Napov and Y. Notay, *An efficient multigrid method for graph Laplacian systems II: robust aggregation*, SIAM J. Sci. Comput. **39**(5), S379–S403, 2017.
- [37] H. Bar-Gera, *Origin-based algorithm for the traffic assignment problem*, Transportation Sci. **36**(4), 398–417, 2002.
- [38] H. Bar-Gera, *Traffic assignment by paired alternative segments (TAPAS)*, Transportation Res. Part B **44**(8–9), 1022–1046, 2010.
- [39] M. Essid and J. Solomon, *Quadratically regularized optimal transport on graphs*, SIAM J. Sci. Comput. **40**(4), A1961–A1986, 2018.